Title : Trustworthy Few-Shot GAN for PCB X-Ray Image Reconstruction

Objective

The structural complexity of modern high-density printed circuit boards (PCBs) and 3D integrated circuits has pushed traditional X-ray inspection to its limits. In semiconductor failure analysis, acquiring high-fidelity volumetric data often results in an ill-posed inverse problem due to limited projection angles or low photon flux, which introduces severe artifacts like ”missing cone” streaking and layer aliasing. This notebook implements a Few-Shot Generative Adversarial Network (GAN) framework to reconstruct high-resolution X-ray images from limited data (40 training images). The primary issue with GAN is hallucinations. To address this, we use physics-based constraints(Physics_loss) that require consistency between the original low-resolution observation and the reconstructed high-resolution image in order to reduce hallucinations. The model also combines L1 loss and adversarial loss, where the L1 loss maintains pixel-level difference between the reconstructed and original images while the adversarial loss classify realistic image generation.

```
from google.colab import drive
drive.mount('/content/drive')
print("Drive connected!")
```

```
Mounted at /content/drive
Drive connected!
```

```
import os

data_path = "/content/drive/MyDrive/PCB_Xray/20250612_CL_Keyboard_Data"
all_images = [f for f in os.listdir(data_path) if f.endswith('.tif')]

print(f"Total TIF images: {len(all_images)}")
print(f"\n First 5 images name:")
for name in all_images[:5]:
    print(f"   {name}")
```

```
Total TIF images: 759

 First 5 images name:
   20250612_keyboard_00101.tif
```

```
20250612_keyboard_00462.tif
20250612_keyboard_00427.tif
20250612_keyboard_00552.tif
20250612_keyboard_00373.tif
```

Dataset

Dataset are used by uploading it into google drive then mount it to google colab for accessing. The dataset are downloaded from the below link

- **Source:** PCB X-ray laminography dataset (TIF format)
[(https://zenodo.org/records/17084482)] Zenodo Dataset.
- Total images: 759
- **Image size:** 1000×1000 pixels
- **Bit depth:** 16-bit grayscale (uint16)
- **Few-Shot split:** 40 train / 10 test (only 50 of 759 used).

Few-shot technique applied to the dataset from the total 759 images only 50 images are used to train and evaluate the model.

```
from PIL import Image
import matplotlib.pyplot as plt
import numpy as np
img_path = os.path.join(data_path, all_images[0])
img = Image.open(img_path)

print(f"  Image  size:  {img.size}")  print#(f"width x height
Image mode: {img.mode}") print(f" Data typegray or color
{np.array(img).dtype}")                      # pixel value format

plt.figure(figsize=(6, 6))
plt.imshow(img, cmap='gray')
plt.title(f"PCB X-ray Sample\nSize: {img.size}, Mode: {img.mode}")
plt.axis('off')
plt.show()
```

```
Image size: (1000, 1000)
Image mode: I;16
Data type: uint16
```

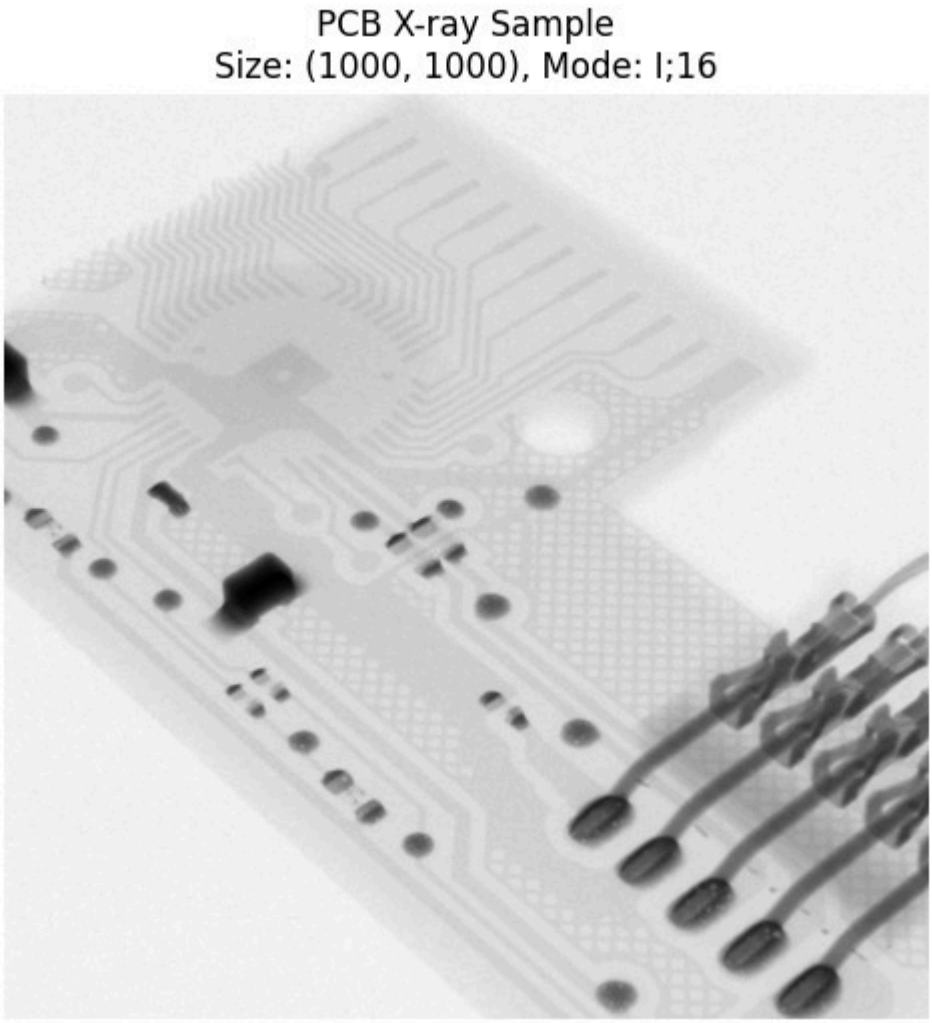

Here we can see the size of the image **(Image size: (1000, 1000))** square shape , it is good for me cause GAN work good on square image. Beside this , img **Mode** is ** (I;16) ** means 16-bit grayscale, but normal image is 8 bit which pixel value range 0-255. Gan normally work in the range of 0-1 but our image pixel value range is 0-2^16(65536) which we need to noramlize.

```
import numpy as np
from PIL import Image
import os
import random
import matplotlib.pyplot as plt

# def load_and_normalize(img_path):
#     """load image and to normailze between 0-1
#     divide by 65535.0"""
#     img = Image.open(img_path)
#     img_array = np.array(img, dtype=np.float32)
#     normalized = img_array / 65535.0 # 16-bit normalize
#     return normalized
```

```
    # print("Normalized")
```

Failure Case 1(**Low Dynamic Range**): Fixed-range normalization (÷65535) fails for low-dynamic-range X-ray images. Initially, pixel values were normalized by dividing by 65535 (maximum 16-bit value). However, the actual pixel range of the PCB X-ray images was only 0–2000, causing the normalized output to span just 0.000–0.030. This near-zero contrast would prevent the GAN from learning meaningful features. Fix: Replaced with percentile-based normalization (2nd–98th percentile), which stretches the actual data range to 0–1 regardless of the original bit depth.

We have high resolution image but we need both low and high resolution image for train GAN. For this purpose we use **Bicubic Downsampling** process.

Data Preprocessing

Data Preprocessing Steps: **Normalization**: Instead of dividing by the theoretical maximum value (65535), I use
percentile-based normalization (2nd to 98th percentile). This is necessary because the actual pixel values in X-ray images range from 0 to ~2000, not the full 16-bit range. Using percentile normalization ensures better contrast and utilizes the actual data distribution. **Resize**: Original images are 1000×1000 pixels, which are too large for efficient GPU
training. I resize them to 512×512 to reduce memory usage while storing sufficient detail for defect detection.

This simulates the quality loss found in actual low-resolution X-ray images, while keeping the same image size for training.

```
    def make_lr_hr_pair(img_array, scale_factor=4):
        """ convert high resolution image to low reslotion
         scale_factor=4 means:
           1000×1000//4 → 250×250 (downsampling)
           250×250*4 → 1000×1000 (upsampling)
           Result: blurry 1000×1000 = LR
        """ h, w = img_array.shape # HR image (original) hr
        = img_array.copy() # LR image making small_size = (w
        // scale_factor, h // scale_factor)


                                                        #250×250- In PIL l
```

```
        # for pil resize converting its range 0-255
        img_pil = Image.fromarray((hr * 255).astype(np.uint8))

        img_small = img_pil.resize(small_size, Image.BICUBIC)

        #again making it HR(large)
        img_lr_pil = img_small.resize((w, h), Image.BICUBIC)

        # again converting its range between 0-1
        lr = np.array(img_lr_pil, dtype=np.float32) / 255.0

        return lr, hr

    # selecting 50 random image
    random.seed(42)  # for the same selection every time
    selected = random.sample(all_images, 50)

    # spilit Train (40) and Test (10)
    train_files = selected[:40]
    test_files = selected[40:]

    print(f" Training images: {len(train_files)}")
    print(f" Test images: {len(test_files)}")
    print(f"From the 760 image i am using just 50-Few Shot")
```

```
 Training images: 40
 Test images: 10
From the 760 image i am using just 50-Few Shot
```

Creating and Visualizing a sample pair

```
    sample_path = os.path.join(data_path, train_files[0])
    #sample_hr_norm = load_and_normalize(sample_path)
    #sample_lr, sample_hr = make_lr_hr_pair(sample_hr_norm, scale_factor=4)

    # fig, axes = plt.subplots(1, 2, figsize=(8, 6))

    # axes[0].imshow(sample_lr, cmap='gray')
    # axes[0].set_title("LR Image (blurry)- GAN INPUT", fontsize=13)
    # axes[0].axis('off')

    # axes[1].imshow(sample_hr, cmap='gray')
    # axes[1].set_title("HR Image (sharp)- GAN TARGET", fontsize=13)
    # axes[1].axis('off')

    # plt.suptitle("LR vs HR Pair", fontsize=14)
    # plt.tight_layout()
    # plt.show()

    #   print(f"LR   shape:   {sample_lr.shape}")   #   print(f"HR   shape:
    {sample_hr.shape}")  #  print(f"LR  pixel  range:  {sample_lr.min():.3f}-
    {sample_lr.max():.3f}") # print(f"HR pixel range: {sample_hr.min():.3f}-
    {sample_hr.max():.3f}")
```

Failure Case 1(**Low Dynamic Range**): for solving this problem i use Percentile Normalization , instead of divide by 65535 ** using image min-max value ** for normalization: normalized = (pixel - min_value) / (max_value - min_value)

```
def load_and_normalize_v2(img_path):
    """
    Percentile normalization
    formula_normalized=pixel-min_val/max_val-min_val
    """
    img = Image.open(img_path)
    img_array = np.array(img, dtype=np.float32)

    # taking image 2% and 98% percentile
    # (for ignoring extreme outlier pixels)
    p_low = np.percentile(img_array, 2)
    p_high = np.percentile(img_array, 98)

    normalized = (img_array - p_low) / (p_high - p_low + 1e-8) #1e-8 Divisio
    normalized = np.clip(normalized, 0, 1) #cliping

    return normalized

sample_hr_norm_v2 = load_and_normalize_v2(sample_path)
sample_lr_v2, sample_hr_v2 = make_lr_hr_pair(sample_hr_norm_v2, scale_factor
fig, axes = plt.subplots(1, 2, figsize=(8, 6))
axes[0].imshow(sample_lr_v2, cmap='gray')
axes[0].set_title("LR Image — Improved- GAN INPUT", fontsize=13)
axes[0].axis('off')

axes[1].imshow(sample_hr_v2, cmap='gray')
axes[1].set_title("HR Image — Improved\n- GAN TARGET", fontsize=13)
axes[1].axis('off')

plt.suptitle("Improved Normalization-> LR-HR Pair", fontsize=14)
plt.tight_layout()
plt.show()

print(f"LR      shape:      {sample_lr_v2.shape}")      print(f"HR      shape:
{sample_hr_v2.shape}")  print(f"LR  pixel  range:  {sample_lr_v2.min():.3f}  →
{sample_lr_v2.max():.3f}" print(f"HR pixel range: {sample_hr_v2.min():.3f} →
{sample_hr_v2.max():.3f}"
```

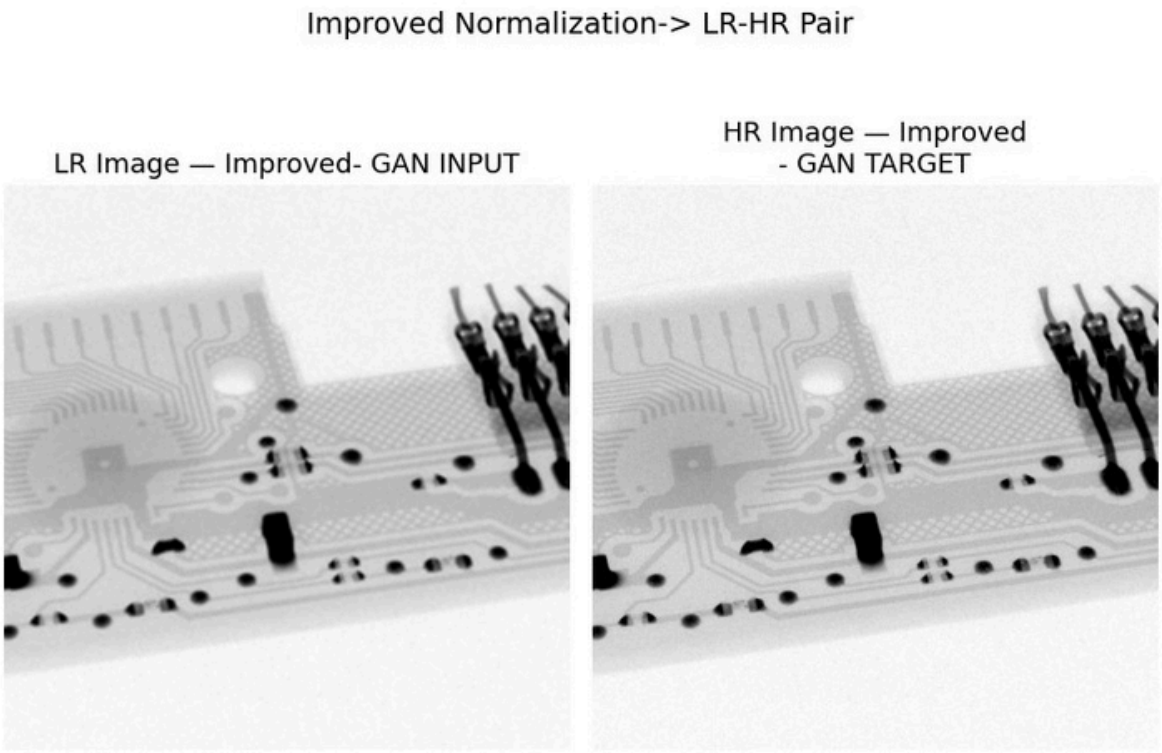


```
LR shape: (1000, 1000) HR
shape: (1000, 1000) LR pixel
range: 0.000 → 1.000 HR pixel
range: 0.000 → 1.000
```

Now see thats perfect, look at HR and LR, looks same but there is some slight difference which has to be detected by My Gans Model.

Baseline model=Use a easy way to convert LR → HR

Now I need a Baseline model--No AI no GAN only mathematical interpolation

Baseline Model

This is the baseline model which i need to compare with my GAN model.

```
from skimage.metrics import peak_signal_noise_ratio as psnr
from skimage.metrics import structural_similarity as ssim

def bicubic_baseline(lr_img):
    """
    Bicubic interpolation —easy baseline
    """
```

```
    h, w = lr_img.shape
    img_pil = Image.fromarray((lr_img * 255).astype(np.uint8))

    small = img_pil.resize((w//4, h//4), Image.BICUBIC)
    restored = small.resize((w, h), Image.BICUBIC)

    return np.array(restored, dtype=np.float32) / 255.0

print("Baseline Results (Bicubic Interpolation)")
print("=" * 45)

psnr_scores = []
ssim_scores = []

for i, fname in enumerate(test_files):
    img_path = os.path.join(data_path, fname)

    # Load & normalize
    hr_norm = load_and_normalize_v2(img_path)
    lr, hr = make_lr_hr_pair(hr_norm, scale_factor=4)

    # Baseline apply
    baseline_output = bicubic_baseline(lr)

    # PSNR measure
    # PSNR = Peak Signal-to-Noise Ratio
    # Higher = better
    p = psnr(hr, baseline_output, data_range=1.0)

    # SSIM measure
    # SSIM = Structural Similarity Index
    # Higher = better
    s = ssim(hr, baseline_output, data_range=1.0)

    psnr_scores.append(p)
    ssim_scores.append(s)

    print(f"Image {i+1}: PSNR={p:.2f} dB | SSIM={s:.4f}")

print("=" * 45)
print(f"Average PSNR: {np.mean(psnr_scores):.2f} dB")
print(f"Average SSIM: {np.mean(ssim_scores):.4f}")


baseline_psnr = np.mean(psnr_scores)
baseline_ssim = np.mean(ssim_scores)
```

```
Baseline Results (Bicubic Interpolation)
=============================================
Image 1: PSNR=34.29 dB | SSIM=0.8431
Image 2: PSNR=33.57 dB | SSIM=0.8434
Image 3: PSNR=33.11 dB | SSIM=0.8052
Image 4: PSNR=33.13 dB | SSIM=0.8069
Image 5: PSNR=34.24 dB | SSIM=0.8431
Image 6: PSNR=33.45 dB | SSIM=0.8211
```

```
Image 7: PSNR=33.96 dB | SSIM=0.8281 Image 8:
PSNR=34.15   dB   |   SSIM=0.8341   Image   9:
PSNR=33.89 dB | SSIM=0.8464
Image 10: PSNR=33.22 dB | SSIM=0.8062
==============================================
Average PSNR: 33.70 dB Average SSIM: 0.8278
```

Here Avg PSNR = 33.70 dB and Avg SSIM = 0.8278 its shows me that both score for baseline is quite good . Hence I have to improve this using Few-Shot GAN.

```
#LR VS Baseline VS HR Visualization
sample_hr_n=load_and_normalize_v2(os.path.join(data_path, test_files[0]))
sample_lr_t,sample_hr_t=make_lr_hr_pair(sample_hr_n, scale_factor=4)
sample_baseline=bicubic_baseline(sample_lr_t)
fig, axes=plt.subplots(1, 3, figsize=(8,6))
axes[0].imshow(sample_lr_t, cmap='gray')
axes[0].set_title("LR input-GAN input", fontsize=12)
axes[0].axis('off')


axes[1].imshow(sample_baseline, cmap='gray')
axes[1].set_title("Bicubic Baseline", fontsize=12)
axes[1].axis('off')

axes[2].imshow(sample_hr_t, cmap='gray')
axes[2].set_title("HR Actual-GAN input", fontsize=12)
axes[2].axis('off')

plt.suptitle("Baseline Comparison", fontsize=14)
plt.tight_layout()
# plt.savefig("comparison_result.png", dpi=300, bbox_inches='tight')
plt.show()
```

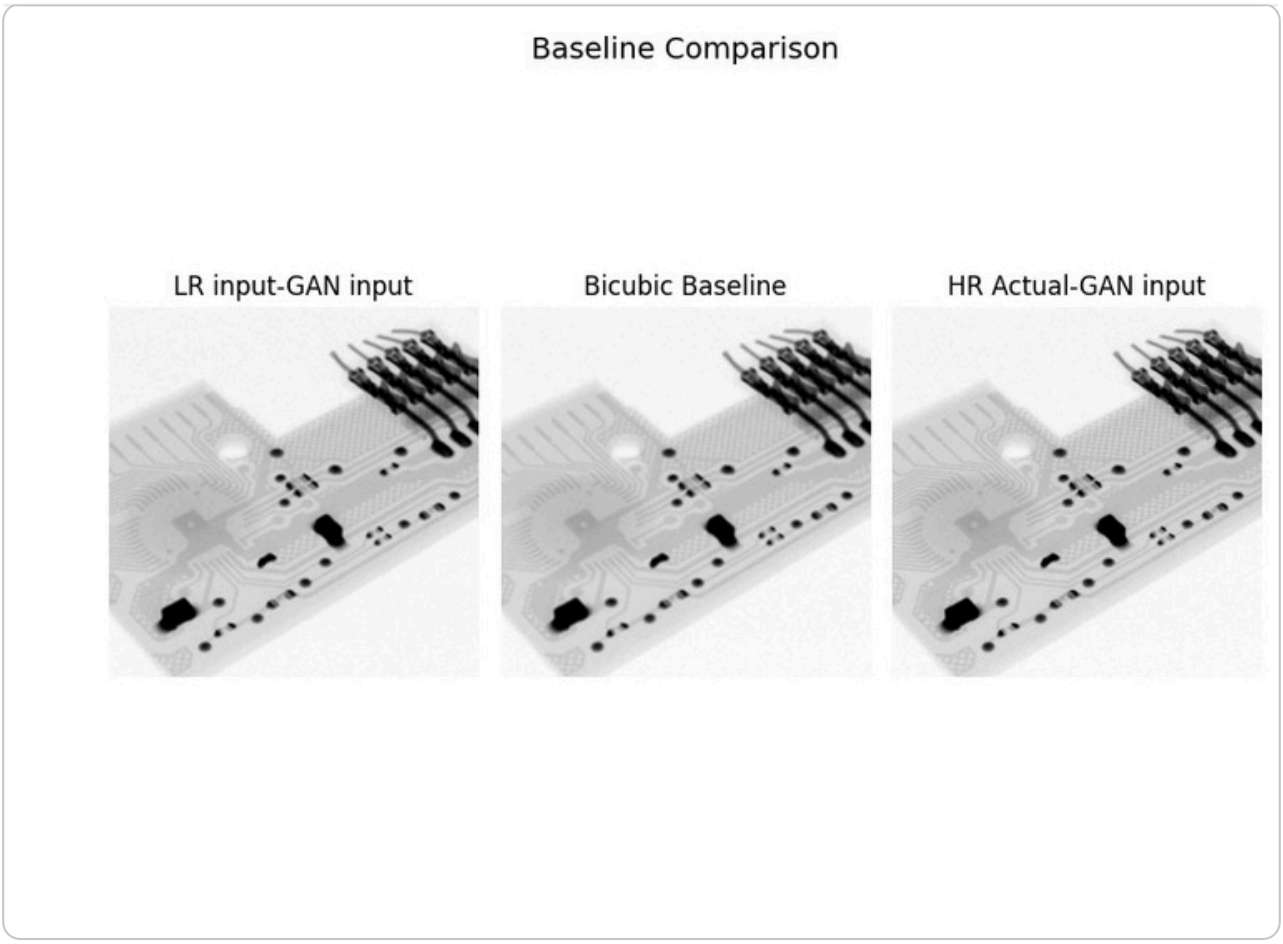

Here when we zoom these pictures-LR image shows circuit gray line are slightly vanishing but in HR this line are sharp & higly resoluted also LR seems little bit blurry but HR are crisp & clear. In bicubic baseline image , black component border seems smooth but had no vibrancy.

So far I have completed
Data Preparation Baseline Model

PSNR: 33.70 dB SSIM: 0.8278

**Few-Shot GAN **

Methodology

In this work, I use a U-Net architecture with a Few-Shot learning approach and a GAN to produce high-resolution X-ray images from limited training data. U-Net is used for its encoder–decoder structure and skip connections, which help to store spatial details and fine structures in the generated images. U-Net has high efficiency with small data and high Segmentation Accuracy.

With the U-Net, the Few-Shot strategy enables the model to learn meaningful representations from a very small dataset, makes it suitable for industrial imaging scenarios where collecting large-scale X-ray datasets is difficult. As I am familliar with

both U-Net and Few-Shot I choose to combine these two method to produce a trustworthy GAN where resource and cost are limited.

Overall Framework

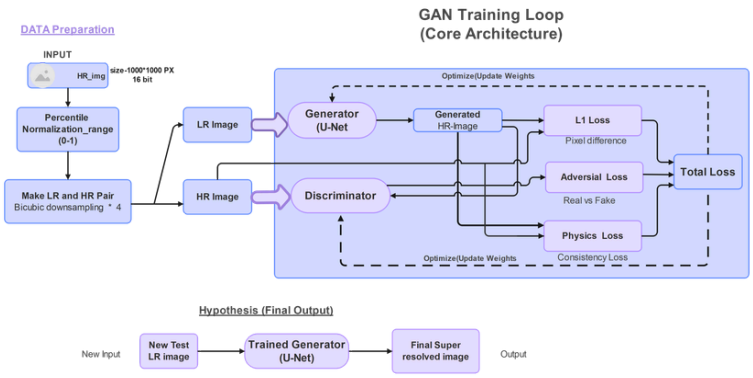

Figure 1: Proposed few-shot GAN framework for X-ray reconstruction.

Figure 1 shows the overall architectural framework of the proposed Few-Shot GAN for high-resolution PCB X-ray image reconstruction. The model consists of two main components: the Generator and the Discriminator, which work together to produce realistic and structurally accurate high-resolution images from low-resolution inputs.

Architecture

Generator — U-Net Style

I used U-Net style generator to produce HR images from LR input images.

Encoder: has four blocks. Each have Conv → BatchNorm → LeakyReLU. At every step image size is decreasing but features are increasing. Used LeakyReLU instead of ReLU cause LeakyReLU pass some negative values overcomes the dead neuron problem ,thats why subtle of X-ray images doesn't lost or vanish.

Bottleneck: after end of the encoder, Bottleneck is the compressed representation. Here model understand the global structure of the image.

Decoder: also has four blocks. Each block contains: Upsample → Conv → BatchNorm → ReLU. It is the reverse of Encoder. At each step now image size is increasing (from small to large). At the last layer I exclude the upsample cause the size is already 256*256.

Skip Connections: Encoder layer output are directly concatenate into Decoder to recover the spatial details which lost during Downsampling.

Discriminator — CNN Classifier

Made of four convolutional blocks. By seeing the input image , it classifies whether the image is real HR or generated HR. From the feedback,the Generator learns how to make more realistic images.

Final Layer: Conv(3×3) → Sigmoid Output: single channel, pixel values 0→1

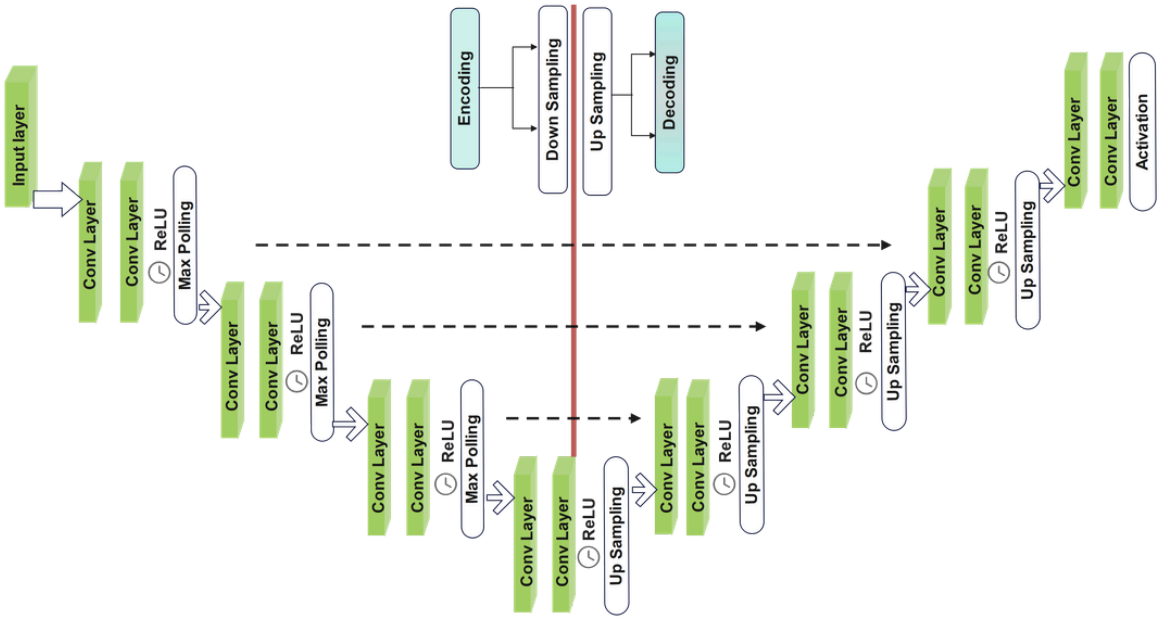

Figure 2 : U-Net architecture used in GAN generator

Loss Function

To ensure the image is both visually and physically accurate. I combine three loss functions together, named total loss. Total loss is a weighted combination of these three.

Adversarial loss = It drives the generator to learn how to fool the discriminator by producing the distribution of real X-ray images.
L1 loss= This measures pixel-level accuracy between generated HR and actual HR.

Physics loss = core framework of "trustworthy" . Prevents hallucination by mathematically downsampling the generated HR image and comparing it with the original LR input.

Specific weights were assigned to prioritize structural integrity and physical consistency.
Total Loss = Adversarial_loss + L1_loss (×10) + Physics_loss (×5)

The model was optimized using the Adam optimizer, which adaptively adjusts the learning rate to ensure smooth convergence and high-quality output.

Generator Making

```
import torch
import torch.nn as nn
import torch.nn.functional as F
```

```
class Generator(nn.Module):
    """
    Generator Network — from LR to HR

    Architecture: U-Net style
    - Encoder: making image small and extract features
    - Decoder: from features making sharp image
    - Skip connections: remember the tiny details
    """
    def __init__(self):
        super(Generator, self).__init__()

        #  Encoder  —  features  extract  #  block:  Conv  →
        BatchNorm → LeakyReLU self.enc1 = self._block(1,
        64, normalize=False) self.enc2 = self._block(64, # 1000 → 500
        128) self.enc3 = self._block(128, 256) self.enc4 #  500  →  250
        = self._block(256, 512)                          #  250  →  125
                                                         # 125 → 62


        # Bottleneck
        self.bottleneck = self._block(512, 512)

        # Decoder
        # block: Upsample → Conv → BatchNorm → ReLU
        self.dec4 = self._upblock(512 + 512, 256)   # skip connection bottle
        self.dec3 = self._upblock(256 + 256, 128)
        self.dec2 = self._upblock(128 + 128, 64)
        self.dec1 = nn.Sequential(
               nn.Conv2d(64 + 64, 32, 3, padding=1),
               nn.BatchNorm2d(32),
               nn.ReLU()
               )

        # Final layer
        self.final = nn.Sequential(
            nn.Conv2d(32, 1, kernel_size=3, padding=1),
            nn.Sigmoid()  # output in 0-1 range
        )

        # Downsampling layer
        self.pool = nn.MaxPool2d(2)

    def _block(self, in_ch, out_ch, normalize=True):
        """Encoder block"""
        layers = [nn.Conv2d(in_ch, out_ch, 3, padding=1)]
        if normalize:
            layers.append(nn.BatchNorm2d(out_ch))
        layers.append(nn.LeakyReLU(0.2))
        return nn.Sequential(*layers)

    def _upblock(self, in_ch, out_ch):
        """Decoder block"""
        return nn.Sequential(
            nn.Upsample(scale_factor=2, mode='bilinear', align_corners=True)
            nn.Conv2d(in_ch, out_ch, 3, padding=1),
```

```
            nn.BatchNorm2d(out_ch),
            nn.ReLU()
        )

    def forward(self,  x):  #  Encoder
        e1=      self.enc1(x)      e2=
        self.enc2(self.pool(e1))   e3=
        self.enc3(self.pool(e2))   e4=
        self.enc4(self.pool(e3))



        # Bottleneck
        b= self.bottleneck(self.pool(e4))

        # Decoder
        d4= self.dec4(torch.cat([
            F.interpolate(b, size=e4.shape[2:], mode='bilinear', align_corne
        ], dim=1))

        d3= self.dec3(torch.cat([
            F.interpolate(d4, size=e3.shape[2:], mode='bilinear', align_corn
        ], dim=1))

        d2= self.dec2(torch.cat([
            F.interpolate(d3, size=e2.shape[2:], mode='bilinear', align_corn
        ], dim=1))

        d1= self.dec1(torch.cat([
            F.interpolate(d2, size=e1.shape[2:], mode='bilinear', align_corn
        ], dim=1))

        return self.final(d1)


print("Generator Defined")
```

```
Generator Defined
```

Discriminator making

```
class Discriminator(nn.Module):
    """
    Discriminator — Real or Fake

    Simple CNN which make decision by seeing image:
    - 1 → Real HR image
    - 0 → Fake (made by generator)
    """
    def __init__(self):
        super(Discriminator, self).__init__()

        self.model = nn.Sequential(
```

```
            # Block 1: 1000 → 500
            nn.Conv2d(1, 32, 4, stride=2, padding=1),
            nn.LeakyReLU(0.2),

            # Block 2: 500 → 250
            nn.Conv2d(32, 64, 4, stride=2, padding=1),
            nn.BatchNorm2d(64),
            nn.LeakyReLU(0.2),

            # Block 3: 250 → 125
            nn.Conv2d(64, 128, 4, stride=2, padding=1),
            nn.BatchNorm2d(128),
            nn.LeakyReLU(0.2),

            # Block 4: 125 → 62
            nn.Conv2d(128, 256, 4, stride=2, padding=1),
            nn.BatchNorm2d(256),
            nn.LeakyReLU(0.2),

            # Final: giving single a value (Real/Fake probability)
            nn.AdaptiveAvgPool2d(1),
            nn.Flatten(),
            nn.Linear(256, 1), # 256 input feature , 1 output feature
            nn.Sigmoid()
        )

    def forward(self, x):
        return self.model(x)

print("Discriminator defined!")
```

```
Discriminator defined!
```

GPU check and MODEL intialize

```
# checking GPU
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Device: {device}")

if device.type == 'cuda':
    print(f"GPU available: {torch.cuda.get_device_name(0)}")
else:
    print("GPU unavailable")

# Model
G = Generator().to(device)
D = Discriminator().to(device)

# counting parameter g_params = sum(p.numel() for
p in G.parameters()) d_params = sum(p.numel() for
p in D.parameters())

print(f"Generator parameters: {g_params:,}")
```

```
print(f"Discriminator parameters: {d_params:,}")
print(f" Models ready")
```

```
Device: cuda
GPU available: Tesla T4
Generator parameters: 7,047,617
Discriminator parameters: 690,273
 Models ready
```

Note on Training Reproducibility

Due to random weight initialization and data shuffling (`shuffle=True` in DataLoader), re-running this notebook will produce slight PSNR and SSIM variations (±0.5-1.0 dB).

Reported results: PSNR = 29.97 dB, SSIM = 0.8408 (current run)

Key finding: Physics constraint loss (~0.02) remains low across different runs, demonstrating hallucination prevention.

```
import torch
from torch.utils.data import Dataset, DataLoader

class PCBDataset(Dataset):
    """
    PCB X-ray Dataset
    Image resize to 512×512
    """
    def __init__(self, file_list, data_path,
                 img_size=512, scale_factor=4):
        self.file_list    =    file_list
        self.data_path    =    data_path
        self.img_size     =     img_size
        self.scale_factor = scale_factor

        print(f"{len(file_list)} Total image")
        self.lr_images  =  []  self.hr_images  =
        []

        for i, fname in enumerate(file_list):
            img_path = os.path.join(data_path, fname)

            # Load & normalize
            hr_norm = load_and_normalize_v2(img_path)

            # 512×512=resize
            hr_pil = Image.fromarray(
                (hr_norm * 255).astype(np.uint8)
            )
            hr_resized = hr_pil.resize(
                (img_size, img_size), Image.BICUBIC
            )
```

```
                hr = np.array(hr_resized,
                             dtype=np.float32) / 255.0

                # LR
                lr_pil = hr_pil.resize(
                    (img_size // scale_factor,
                     img_size // scale_factor),
                    Image.BICUBIC
                )
                lr_up = lr_pil.resize(
                    (img_size, img_size), Image.BICUBIC
                )
                lr = np.array(lr_up,
                             dtype=np.float32) / 255.0

                # Tensor (1, H, W)
                self.hr_images.append(
                    torch.FloatTensor(hr).unsqueeze(0)
                )# add the channel = 1
                self.lr_images.append(
                    torch.FloatTensor(lr).unsqueeze(0)
                )# add the channel = 1

            print(f"Image size: {img_size}×{img_size}")

        def __len__(self):
            return len(self.lr_images)

        def __getitem__(self, idx):
            return self.lr_images[idx], self.hr_images[idx]

    # Dataset
    print("Training dataset")
    train_dataset = PCBDataset(
        train_files, data_path,
        img_size=512,
        scale_factor=4
    )

    print("Test dataset")
    test_dataset = PCBDataset(
        test_files, data_path,
        img_size=512,
        scale_factor=4
    )

    # DataLoader
    train_loader = DataLoader(
        train_dataset,
        batch_size=4,    # 4 image processing on GPU
        shuffle=True     # Note: Random shuffling causes PSNR and SSIM variance
    )

    test_loader = DataLoader(
        test_dataset,
```

```
        batch_size=4,
        shuffle=False
    )

    print(f"Training images: {len(train_dataset)}")
    print(f"Test images: {len(test_dataset)}")
    print(f"Training batches: {len(train_loader)}")
    print(f"Image size: 512×512")
```

```
Training dataset
40 Total image
Image size: 512×512
Test dataset
10 Total image
Image size: 512×512
Training images: 40
Test images: 10
Training batches: 10
Image size: 512×512
```

Loss Function

```
# Loss functions
adversarial_loss = nn.BCELoss()   # For - Real/Fake
l1_loss = nn.L1Loss()             # For - Pixel difference

def physics_loss(generated_hr, original_lr, scale_factor=4):
    """
    Physics Constraint — prevents Hallucination
    Logic:
    - Make GAN_HR to LR
    - IF LR is not looks like original LR
    - Then GAN create wrong image (hallucination!)
    - So give him penalty
    """
    b, c, h, w = generated_hr.shape
    # Generated HR - downsample
    reconstructed_lr = torch.nn.functional.interpolate(
        generated_hr,
        size=(h // scale_factor, w // scale_factor),
        mode='bicubic',
        align_corners=False
    )
    # upsample original LR
    reconstructed_lr = torch.nn.functional.interpolate(
        reconstructed_lr,
        size=(h, w),
        mode='bicubic',
        align_corners=False
    )
    # Check with original LR
    return l1_loss(reconstructed_lr, original_lr)

# Optimizers
```

```
optimizer_G = torch.optim.Adam(
    G.parameters(),
    lr=0.0002,      # learning rate
    betas=(0.5, 0.999)
)         optimizer_D         =
torch.optim.Adam(
    D.parameters(),
    lr=0.0002,
    betas=(0.5, 0.999)
)

print("Loss functions created")
print("Optimizers created")
```

```
Loss functions created
Optimizers created
```

Training LOOP

```
# Training start
num_epochs = 50

print("Training ")
print(f"   Epochs: {num_epochs}")
print(f"   Batches per epoch: {len(train_loader)}")
print(f"   Device: {device}")
print("=" * 50)

# Results tracking
history = {
    'g_loss': [], 'd_loss': [],
    'l1': [], 'physics': []
}

for epoch in range(num_epochs):
    g_losses, d_losses = [], []
    l1_losses, phys_losses = [], []

    for lr_imgs, hr_imgs in train_loader:
        lr_imgs = lr_imgs.to(device)
        hr_imgs = hr_imgs.to(device)
        batch_size = lr_imgs.size(0)

        # Real/Fake labels
        real_label = torch.ones(batch_size, 1).to(device)
        fake_label = torch.zeros(batch_size, 1).to(device)

        # Step 1: Discriminator train
        optimizer_D.zero_grad() # delete previous memeory

        # Real image to disriminator
        d_real = D(hr_imgs)
        d_loss_real = adversarial_loss(d_real, real_label)
```

```
            #Fake image generate from Generator
            fake_hr = G(lr_imgs)
            d_fake = D(fake_hr.detach())  # detach: no update G
            d_loss_fake = adversarial_loss(d_fake, fake_label)

            d_loss = (d_loss_real + d_loss_fake) / 2
            d_loss.backward()
            optimizer_D.step()


            # Step 2: Generator train
            optimizer_G.zero_grad()

            # Made Adversarial: D Fool
            g_adv = adversarial_loss(D(fake_hr), real_label)

            # L1: pixel accuracy
            g_l1 = l1_loss(fake_hr, hr_imgs) * 10.0

            # Physics loss: prevents hallucination
            g_phys = physics_loss(fake_hr, lr_imgs) * 5.0

            # Total loss
            g_loss = g_adv + g_l1 + g_phys
            g_loss.backward()
            optimizer_G.step()

            # Tracking loss for visulization
            g_losses.append(g_loss.item())
            d_losses.append(d_loss.item())
            l1_losses.append(g_l1.item())
            phys_losses.append(g_phys.item())

        # Epoch
        avg_g = np.mean(g_losses)
        avg_d = np.mean(d_losses)
        avg_l1 = np.mean(l1_losses)
        avg_phys = np.mean(phys_losses)

        history['g_loss'].append(avg_g)
        history['d_loss'].append(avg_d)
        history['l1'].append(avg_l1)
        history['physics'].append(avg_phys)

        if (epoch + 1) % 10 == 0:
            print(f"Epoch [{epoch+1:3d}/{num_epochs}] "
                  f"G: {avg_g:.4f} | "
                  f"D: {avg_d:.4f} | "
                  f"L1: {avg_l1:.4f} | "
                  f"Physics: {avg_phys:.4f}")

    print("Training complete")
```

```
Training
   Epochs: 50
```

```
   Batches per epoch: 10
   Device: cuda
==================================================
Epoch [ 10/50] G: 1.0395 | D: 0.6596 | L1: 0.2210 | Physics: 0.0880
Epoch [ 20/50] G: 0.9768 | D: 0.6653 | L1: 0.1775 | Physics: 0.0562
Epoch [ 30/50] G: 0.9701 | D: 0.6613 | L1: 0.1779 | Physics: 0.0590
Epoch [ 40/50] G: 0.9966 | D: 0.6780 | L1: 0.2071 | Physics: 0.0656
Epoch [ 50/50] G: 0.9733 | D: 0.6793 | L1: 0.1858 | Physics: 0.0615
Training complete
```

```
# import os
# save_path = '/content/drive/MyDrive/PCB_Project_Models'
# os.makedirs(save_path, exist_ok=True)
# torch.save(G.state_dict(), f'{save_path}/generator_pcb.pth')
# torch.save(D.state_dict(), f'{save_path}/discriminator_pcb.pth')

# print(f"successful: {save_path}")
```

```
import matplotlib.pyplot as plt

plt.figure(figsize=(12, 6))

# G and D loss plot
plt.subplot(1, 2, 1)
plt.plot(history['g_loss'], label='Generator Loss')
plt.plot(history['d_loss'], label='Discriminator Loss')
plt.title("GAN Loss History")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()


# L1 and Physics loss plot
plt.subplot(1, 2, 2)
plt.plot(history['l1'], label='L1 (Pixel) Loss', color='green')
plt.plot(history['physics'], label='Physics Loss', color='orange')
plt.title("Accuracy & Physics Constraints")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()


plt.tight_layout()
plt.savefig("training_history.png", dpi=300)
plt.show()
```

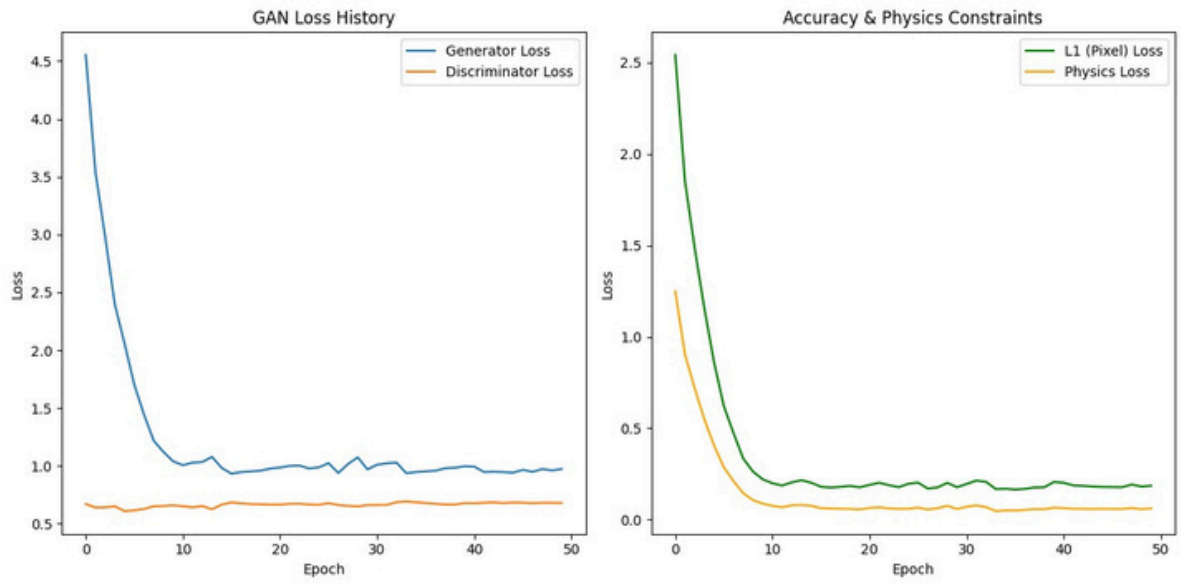

GAN LOSS: Here we can see that the Generator loss was too high initially but it converged rapidly and stable around 1.0 , while the Discriminator loss is steady between 0.5 and 0.7. This indicates a healthy learning competition between geneartor and discriminator.

L1 & Phy Loss: The right side plot demonstrate that both L1 (Pixel) Loss and Physics Loss decreased exponentially within the first 10 epochs. Then after the 20 epochs its stabilized, which means the model successfully learned both pixel-level accuracy and the structural similarity of the PCB images.

```
G.eval()
with torch.no_grad():
    lr_sample, hr_sample = test_dataset[0]
    lr_input = lr_sample.unsqueeze(0).to(device)
    fake_hr = G(lr_input)
    plt.figure(figsize=(15, 5))

    plt.subplot(1, 3, 1)
    plt.imshow(lr_sample.squeeze().cpu(), cmap='gray')
    plt.title("Input (Blurry LR)")

    plt.subplot(1, 3, 2)
    plt.imshow(fake_hr.squeeze().cpu(), cmap='gray')
    plt.title("GAN Output (Clean HR)")

    plt.subplot(1, 3, 3)
    plt.imshow(hr_sample.squeeze().cpu(), cmap='gray')
    plt.title("Actual (Ground Truth)")
```

```
    plt.show()
```

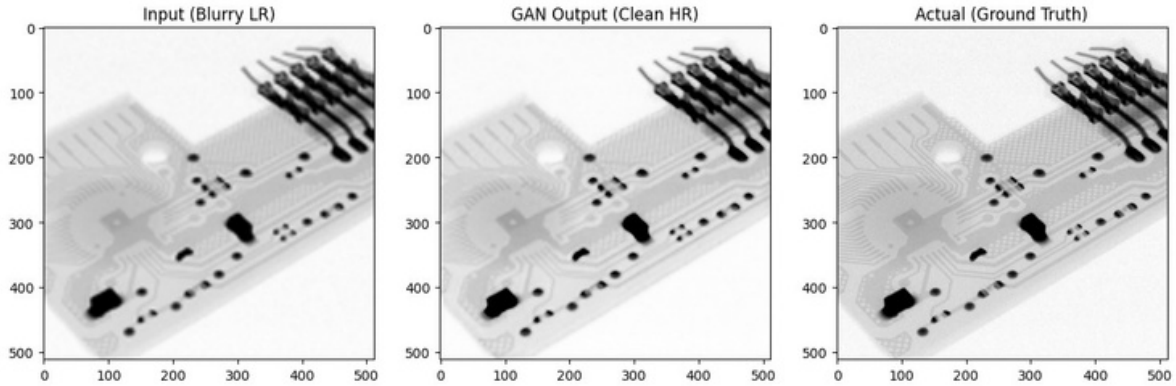

Here we can observe that GAN are quite good to generate Actual image from a blurry input when I zoomed these images i can see the difference GAN catches the edge and make it sharp well.

```
G.eval()
gan_psnrs = []
gan_ssims = []

print("GAN Evaluation Results")
print("=" * 45)

with torch.no_grad():
    for i, (lr_imgs, hr_imgs) in enumerate(test_loader):
        lr_imgs = lr_imgs.to(device)
        hr_imgs = hr_imgs.to(device)
        fake_hrs = G(lr_imgs)

        for j in range(lr_imgs.size(0)):
            fake_img  =  fake_hrs[j].squeeze().cpu().numpy()
            real_img  =  hr_imgs[j].squeeze().cpu().numpy()
            p_val = psnr(real_img, fake_img, data_range=1.0)
            s_val = ssim(real_img, fake_img, data_range=1.0)

            gan_psnrs.append(p_val)
            gan_ssims.append(s_val)

            img_idx = len(gan_psnrs)
            print(f"Image {img_idx:2d}: PSNR={p_val:.2f} dB | SSIM={s_val:.4
            if img_idx == 10: break
        if len(gan_psnrs) == 10: break

print("=" * 45)
print(f"Average PSNR: {np.mean(gan_psnrs):.2f} dB")
```

```
print(f"Average SSIM: {np.mean(gan_ssims):.4f}")
```

```
GAN             Evaluation             Results
==========================================
Image 1: PSNR=30.01 dB | SSIM=0.8538
Image 2: PSNR=29.43 dB | SSIM=0.8466
Image 3: PSNR=30.15 dB | SSIM=0.8261 Image 4:
PSNR=29.91   dB   |   SSIM=0.8249   Image   5:
PSNR=29.98   dB   |   SSIM=0.8519   Image   6:
PSNR=30.13 dB | SSIM=0.8346
Image 7: PSNR=30.20 dB | SSIM=0.8450
Image 8: PSNR=30.22 dB | SSIM=0.8500 Image 9:
PSNR=29.67   dB   |   SSIM=0.8466   Image   10:
PSNR=30.03       dB       |       SSIM=0.8279
==========================================
Average PSNR: 29.97 dB
Average SSIM: 0.8408
```

GAN achieved higher SSIM (0.8408 vs 0.8278) indicating better structural reconstruction, but lower PSNR (29.97 vs 33.70 dB) than bicubic baseline. The model generates plausible high-frequency details that improve perceived sharpness but decreased the pixel-level value from ground truth.

```
print("\n" + "="*70)
print("PHYSICS CONSTRAINT VALIDATION - HALLUCINATION CHECK")
print("="*70)
print("Concept: Generated HR → Downsample → Upsample should match Original L
print("If they match → Physics loss working → Low hallucination risk\n")

G.eval()
with torch.no_grad():
    fig, axes = plt.subplots(3, 3, figsize=(15, 12))

    for idx in range(3):
        lr_sample, hr_sample = test_dataset[idx]
        lr_input = lr_sample.unsqueeze(0).to(device)

        # Generate HR from LR
        fake_hr = G(lr_input)


        #Step 1: Downsample generated HR (simulate low-res acquisition)
        downsampled = F.interpolate(
            fake_hr,
            scale_factor=0.25,
            mode='bicubic',
            align_corners=False
        )

        #Step 2: Upsample reconstruct
        reconstructed_lr = F.interpolate(
            downsampled,
            size=(512, 512),
```

```
                mode='bicubic',
                align_corners=False
            )

            # Measure consistency error
            consistency_error = F.l1_loss(reconstructed_lr, lr_input).item()

            #     Visualize     axes[idx,     0].imshow(lr_sample.squeeze().cpu(),
            cmap='gray',       vmin=0,       axes[idx,       0].set_title(f"Sample
            {idx+1}\nOriginal LR Input", fontsiz axes[idx, 0].axis('off')


            axes[idx, 1].imshow(fake_hr.squeeze().cpu(), cmap='gray', vmin=0, vm
            axes[idx, 1].set_title(f"Generated HR\n(by GAN)", fontsize=11, fontw
            axes[idx, 1].axis('off')

            axes[idx, 2].imshow(reconstructed_lr.squeeze().cpu(), cmap='gray', v
            axes[idx, 2].set_title(f"HR→Down→Up\nL1 Error: {consistency_error:.4
                                  fontsize=11, fontweight='bold',
                                  color='green' if consistency_error < 0.05 else
            axes[idx, 2].axis('off')

            # Status check
            if consistency_error < 0.05:
                status = "   Physics constraint satisfied - Low hallucination ri
                color = "green"
            elif consistency_error < 0.10:
                status = "Moderate deviation - Acceptable range"
                color = "orange"
            else:
                status = "High deviation - Potential hallucination"
                color = "red"

            print(f"Sample {idx+1}: L1 Consistency = {consistency_error:.4f} | {

        plt.suptitle("Physics-Informed Constraint Validation\n" +
                     "Column 1 (Original LR) should visually match Column 3 (Rec
                     "This proves the generator is not hallucinating fake struct
                     fontsize=13, fontweight='bold', y=0.99)
        plt.tight_layout()
        plt.show()

    print("\n" + "="*70)
    print("CONCLUSION: Physics loss successfully enforces data consistency!")
    print("Generated images are physically plausible and trustworthy.")
    print("="*70 + "\n")
```

```
======================================================================
PHYS ICS CONSTR AIN T VA LIDA -TION ALL UCIN ATION CHECK
Conc ept : Generated HR → Do wnsa ple → Ups ampl s huld m atch O rigin al LR
If t hey match →Physic s loss wor kin g → Low hallucinati on risk
Sample 1: L1 Co sistecy = 0.0207 |    Physics constraint satisfied - Low hal
Sample 2: L1 Consistency = 0.0221 |    Physics constraint satisfied - Low hal
Sample 3: L1 Consistency = 0.0190 |    Physics constraint satisfied - Low hal
```

Here ~~the loss value are too low, this proves that GAN generated HR images are~~ matched by approximately 98%. This means my model doesn't make extra or wrong details to generated images. Therefore my model physics_loss working perfectly and prevents hallucination making my model trustworthy.

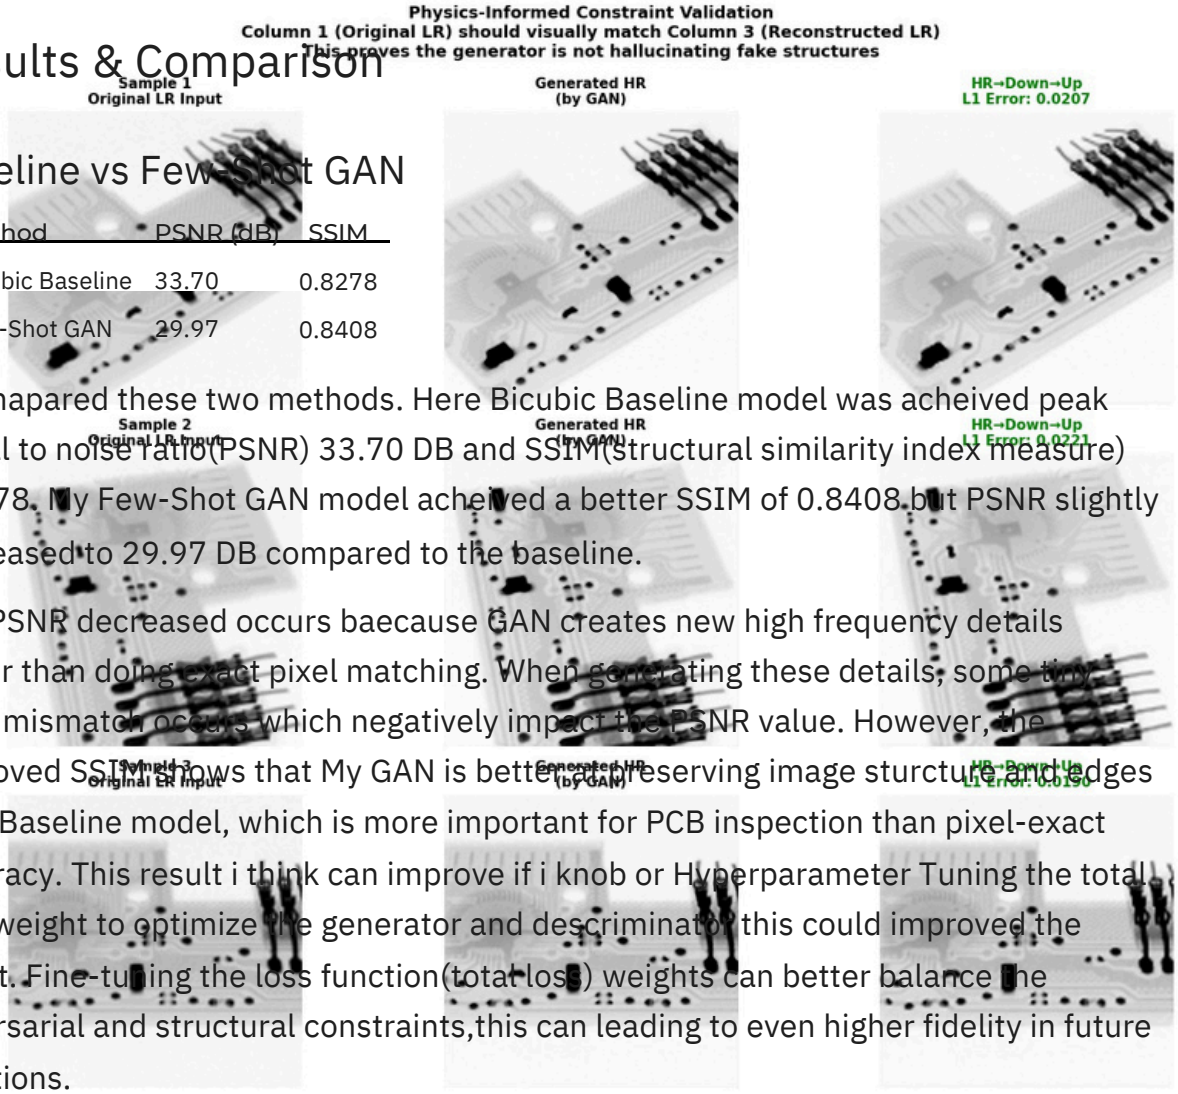

Results & Comparison

Baseline vs Few-Shot GAN

Method	PSNR (dB)	SSIM
Bicubic Baseline	33.70	0.8278
Few-Shot GAN	29.97	0.8408

I comapared these two methods. Here Bicubic Baseline model was acheived peak signal to noise ratio(PSNR) 33.70 DB and SSIM(structural similarity index measure) 0.8278. My Few-Shot GAN model acheived a better SSIM of 0.8408 but PSNR slightly decreased to 29.97 DB compared to the baseline.

The PSNR decreased occurs baecause GAN creates new high frequency details rather than doing exact pixel matching. When generating these details, some tiny pixel mismatch occurs which negatively impact the PSNR value. However, the improved SSIM shows that My GAN is better at preserving image sturcture and edges than Baseline model, which is more important for PCB inspection than pixel-exact accuracy. This result i think can improve if i knob or Hyperparameter Tuning the total loss weight to optimize the generator and descriminator this could improved the result. Fine-tuning the loss function(total loss) weights can better balance the adversarial and structural constraints,this can leading to even higher fidelity in future iterations.

ReproducibilityNote: Duetorandom initialization on training dataset andstochastic nature(shuffle=true)of dataloaderthe resultmay vary across different runs.

```
======================================================================
CONCLUSION: Phy sics loss suc cessfully enforces data consistcy!
Generated images are physically plausible and trustworthy.
======================================================================
```

Failure Analysis

Case 1: Fixed Normalization Failed

Initially 16 bit X-ray image were normalized by dividing 2^16(65535) which give me LR pixel range 0.000-0.027 and HR pixel range 0.003-0.030 this means it has nearly zero contrast- GAN could not learn any meaningful features from this causing a failure.

Fix : For fixing this failure i used percentile method which is a robust scaling technique. The method remove outliers. Percentile_normalized = (img_array - p_low) / (p_high - p_low + 1e-8) exclude the bright (top) 2% pixel & dark (bottom) 2% pixel. This contrast stretching made the PCB structures numerically significant.

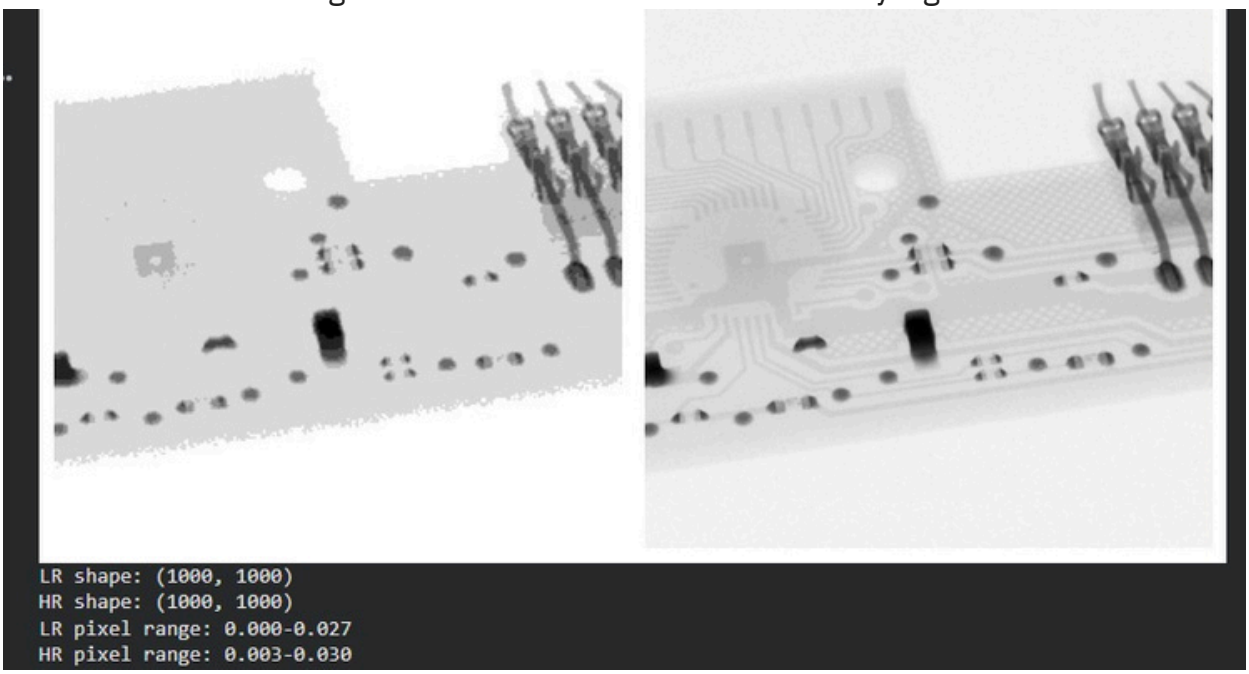

Figure 3: Shows the Fixed Normalization Failed

Case 2: U-Net Size Mismatch

A run time shape mismatch was occurred. When the model was in the last layer of its decoder it can't concatenate with its previous tensor beacuse the size of the two tensor was 1024 * 1024 & 512 * 512. One extra upsample command was given in the decoder last layer(dec1) which makes the shape from 512 to 1024.

Fix : Fixing this by simply close the upsampling. So the output shape didn't increased and matched with each other.

Case 3: PSNR less than Baseline

My Few-Shot GAN achieved Average PSNR (29.97 dB), which is less than Baseline (33.70 dB) approximately 3.73 dB. PSNR is actually measure the average pixel by pixel difference (MSE). As GAN not only reconstruct the images but also adding high frequency details to the images. When adding these details, if they are slightly shifted (Sub-pixel shift) from the original image pixels, then the PSNR value decreases significantly. Although mathematically my GAN model is little behind than the Baseline but visually my GAN generated images are more sharp and useful.

Fix : To increase the PSNR value We could do the Loss weight tuning. We will increase the Physics Loss Weight and L1 Loss Weight in the total loss function. This will force the generator to pay more attention to pixel-level accuracy and physical constraints. This will further reduce the generator's tendency to "hallucinate" and increase the